



Writing a Windows Kernel Driver in an Hour (or so)

Pavel Yosifovich
@zodiacon
zodiacon@live.com

About Me

- Developer, Trainer, Author, Speaker
- Book author
 - "Windows Internals 7th edition, Part 1" (co-author, 2017)
 - "Windows 10 System Programming, Part 1" (2020)
 - "Windows 10 System Programming, Part 2" (2021)
 - "Windows Kernel Programming, 2nd ed." (2023)
- *Pluralsight* and *Pentester Academy* course author
- Author of several open-source tools (<http://github.com/zodiacon>)
- Website: <http://scorpiosoftware.net>



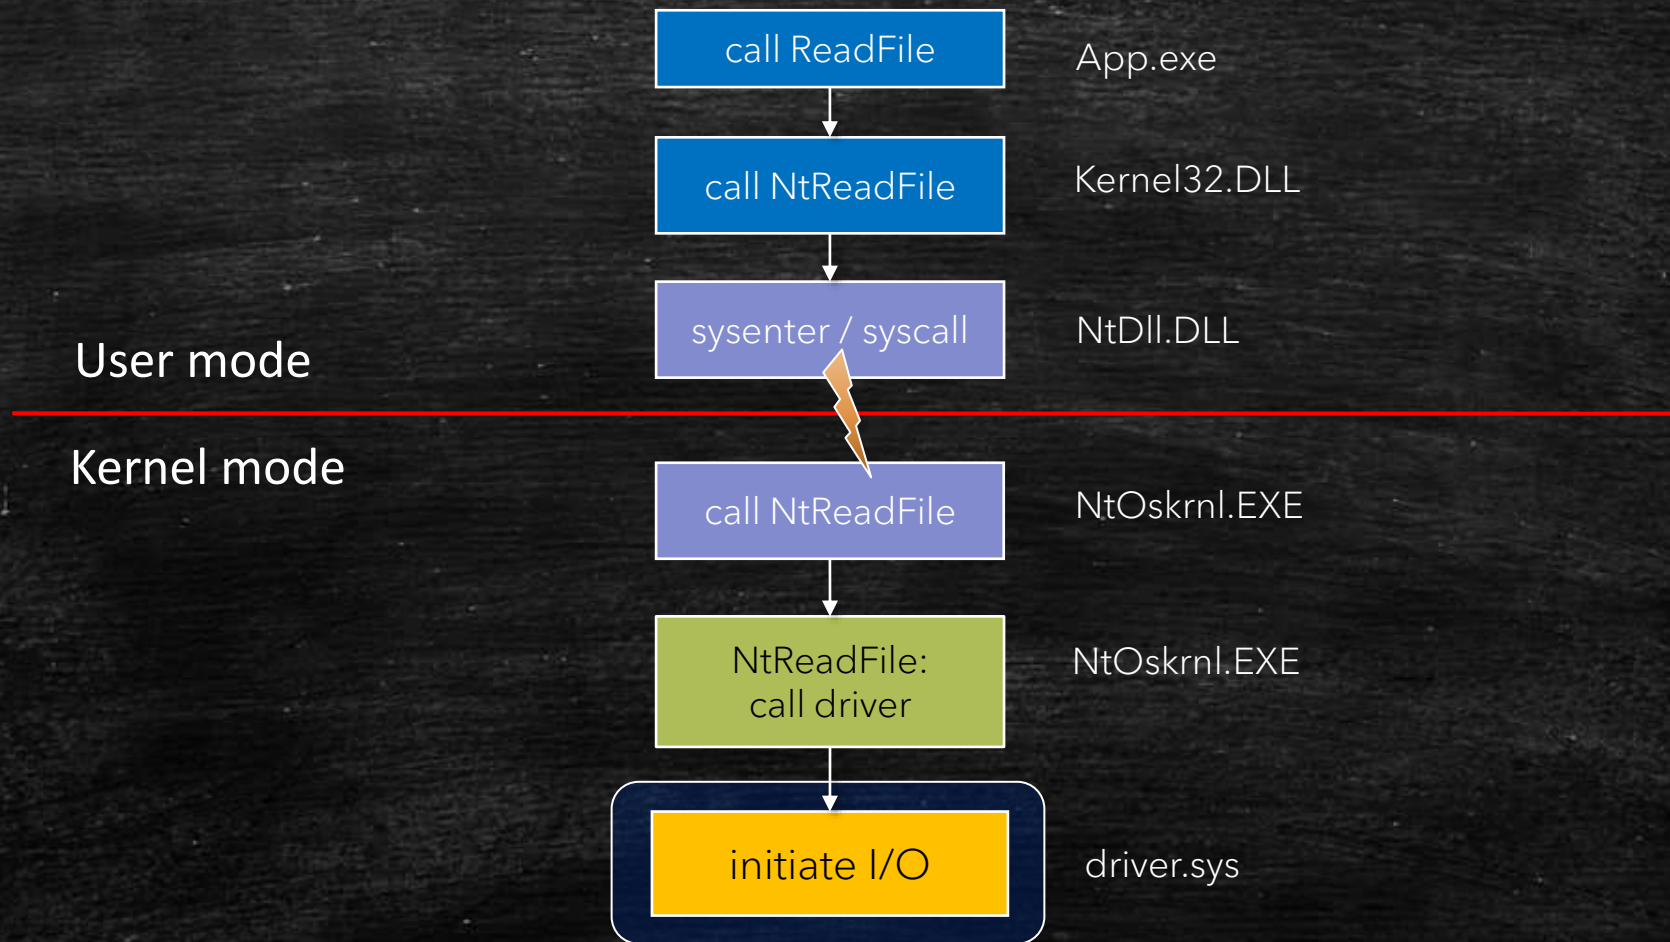
Overview

- Kernel Driver Basics
- Invoking Drivers
- **DriverEntry**
- Drivers, Devices, and Symbolic Links
- Handling Requests
- Testing and Debugging

Kernel Device Drivers

- Always execute in kernel mode
 - Use the kernel mode stack of a thread
 - Image part of system space
 - Unhandled exceptions will crash the system
- Typically has a SYS file extension
- Usually invoked by client code
 - e.g. `ReadFile`, `WriteFile`, `DeviceIoControl`
- Exports entry points for various functions
 - Called by system code when appropriate
- System handles all device independent aspects of I/O
 - No need for hardware specific code or assembly

Invoking a Driver



Accessing Devices

- A client that wants to communicate with a device, must open a handle to the device
 - **CreateFile** or **CreateFile2** from user mode
 - **ZwCreateFile** or **ZwOpenFile** from kernel mode
- **CreateFile** accepts a "filename" which is actually a device symbolic link
 - "file" being just one specific case
 - For devices, the name should have the format \\.**name**
 - Cannot access non-local device

The CreateFile API

- Creates a file object
 - Creation flags must specify **OPEN_EXISTING** for devices
 - Driver sees **CreateFile** as an IRP with major code **IRP_MJ_CREATE**
 - Returns a handle to the file object
 - Returns **INVALID_HANDLE_VALUE** (-1) if fails
 - Call **GetLastError** to get reason for failure

```
HANDLE CreateFile(  
    _In_          LPCTSTR lpFileName,           // name of file or device  
    _In_          DWORD dwDesiredAccess,        // access mode (read, write, etc.)  
    _In_          DWORD dwShareMode,           // share mode (for files)  
    _In_opt_ LPSECURITY_ATTRIBUTES sa,         // security descriptor and handle inheritance  
    _In_          DWORD dwCreationDisposition, // creation flags (OPEN_EXISTING for devices)  
    _In_          DWORD dwFlagsAndAttributes,  // more flags  
    _In_opt_      HANDLE hTemplateFile);       // template file to copy attributes from
```

Setting Up For Kernel Development

- Install Visual Studio 2022 (any edition)
- Install the latest Windows 11 Software Development kit
- Install the latest Windows 11 Driver Kit (WDK)
 - The latest WDK allows building drivers for Windows 10 and later
 - The WDK installs a bunch of project templates for driver development
- Optionally, install the new WinDbg

The Kernel API

- Most of the kernel code is in *NtOsKrnl.exe*
- Some implementation is in *Hal.dll*
- The WDK documents about third of exported functions
- Most functions have a prefix indicating the component they are implemented in
- Main header file: *<ntddk.h>*

Some Kernel API Prefixes

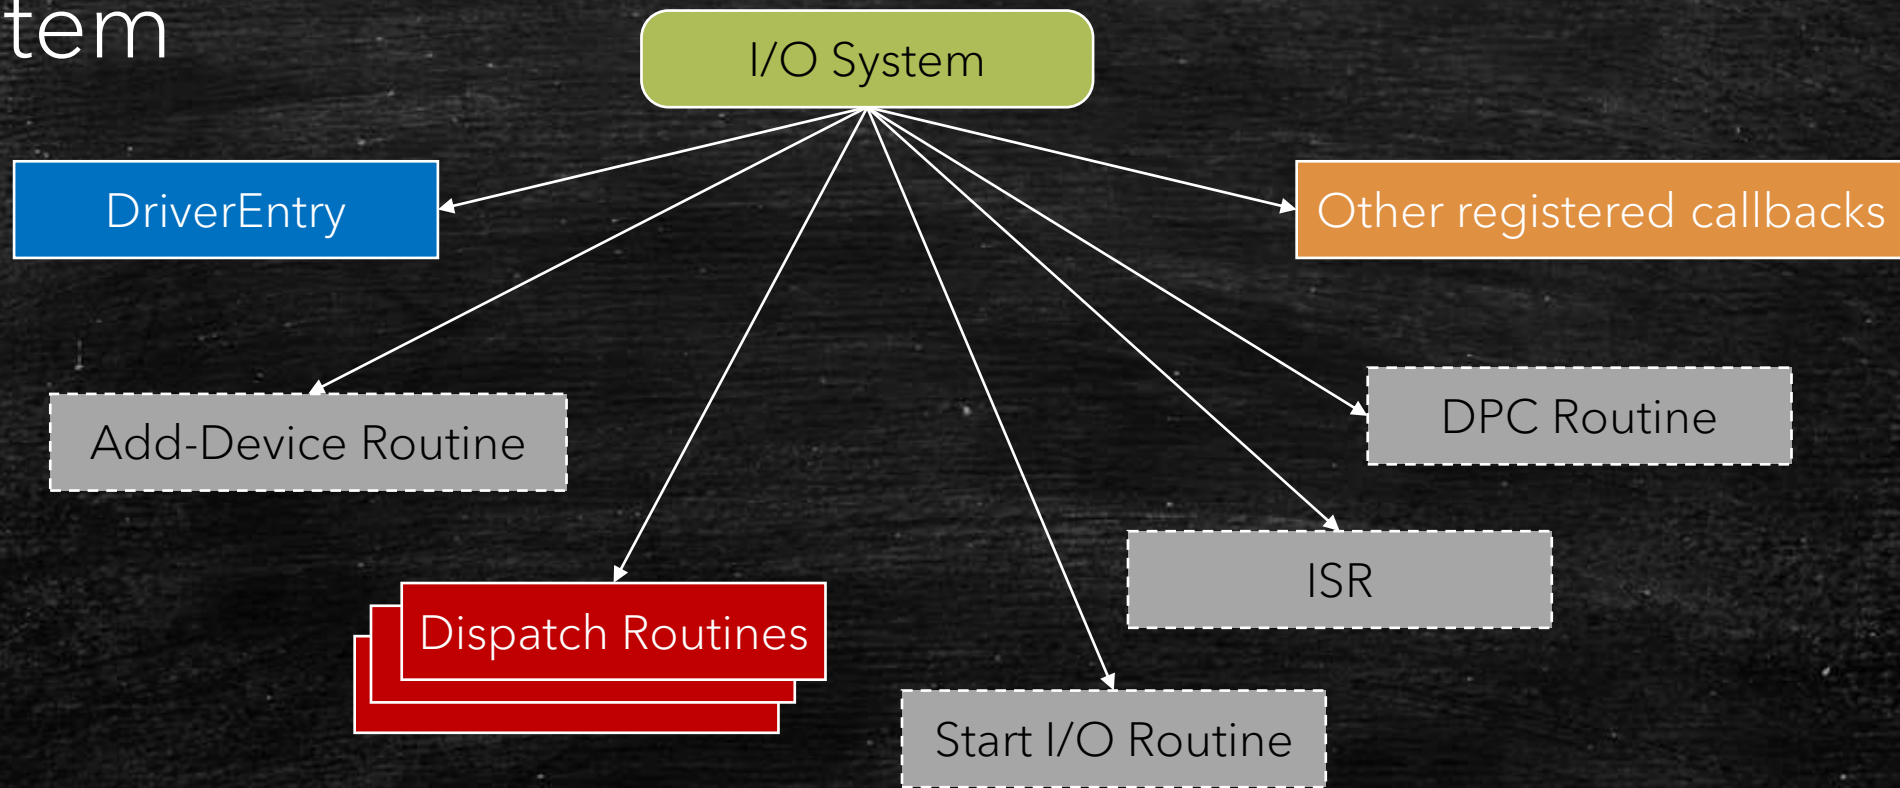
Ex	- General executive functions
Ke	- General kernel functions
Cc	- Cache Controller (Manager)
Mm	- Memory Manager
Rtl	- General runtime library
FsRtl	- File System Runtime Library
Flt	- File system mini-filters
Ob	- Object Manager
Io	- I/O Manager
Se	- Security
Ps	- Process Structure
Po	- Power Manager
Wmi	- Windows Management Instrumentation
Zw	- Native API wrappers
Hal	- Hardware Abstraction Layer

Functions and Error Codes

- Most kernel functions return **NTSTATUS**
 - 32-bit integer
 - MSB=1 indicates error
- Many status values
 - **STATUS_SUCCESS** = 0 (standard success)
 - Many failure status values
 - **STATUS_UNSUCCESSFUL** being the most generic
- If returned to user mode, turned into **ERROR_XXX**
- Check success with **NT_SUCCESS**(status) macro

Anatomy of a Driver

- A driver exports functionality, callable by the I/O system



Strings

- Most kernel APIs work with **UNICODE_STRING** structures
 - **Length** and **MaximumLength** are in bytes
 - Need *not* be NULL terminated
- Useful functions
 - **RtlInitUnicodeString**
 - Initializes a pre-allocated string (e.g. constant string)
 - **RtlCopyUnicodeString**
 - Copy from an existing **UNICODE_STRING**
 - Destination string must be initialized
 - **RtlCompareUnicodeString**
 - Case sensitive or insensitive comparison
 - **RTL_CONSTANT_STRING** macro to initialize with a literal string

```
typedef struct _UNICODE_STRING {  
    USHORT Length;  
    USHORT MaximumLength;  
    PWCH Buffer;  
} UNICODE_STRING;  
typedef UNICODE_STRING *PUNICODE_STRING;
```

The DriverEntry Function

- Must have for any kernel driver

```
NTSTATUS DriverEntry(  
    _Inout_ PDRIVER_OBJECT DriverObject,  
    _In_ PUNICODE_STRING RegistryPath);
```


- Called with the driver object partially initialized using a system worker thread at IRQL 0
- Registry key points to the "Software" key
 - **HKLM\System\CurrentControlSet\Services\ServiceName**
 - String is valid in **DriverEntry** only (deep copy if needed in other functions)
- Returning anything other than **STATUS_SUCCESS** causes driver to unload
 - Unload routine *not* called

Initializing the Driver Object

- Set the Unload routine
 - **DriverUnload** member of **DRIVER_OBJECT**
- Set supported major functions
 - **MajorFunction** array of function pointers
 - Unset major functions are initialized to "unsupported function"

Driver and Device Objects

- Drivers are represented in memory using a **DRIVER_OBJECT** structure
 - Created by the I/O system
 - Provided to the driver in the **DriverEntry** function
 - Holds all exported functions
- Device objects are created by the driver on a per-device basis
 - Within the **AddDevice** driver callback (for P&P drivers) or in **DriverEntry** (for software drivers)
 - Represented by the **DEVICE_OBJECT** structure
- I/O system is device-centric, not driver-centric



Exporting Entry Points

- **DriverEntry**'s main job is to export entry points for its main functions
 - **DriverUnload**
 - The unload routine, to be called when driver is about to be unloaded from memory
 - **MajorFunction**
 - An array of function pointers for the various types of IRPs the driver is going to handle

Creating a Device Object

- Define a device extension structure to hold all device-specific data
- Call **IoCreateDevice**

```
NTSTATUS IoCreateDevice(  
    _In_      PDRIVER_OBJECT DriverObject,  
    _In_      ULONG DeviceExtensionSize,  
    _In_opt_  PUNICODE_STRING DeviceName,  
    _In_      DEVICE_TYPE DeviceType,  
    _In_      ULONG DeviceCharacteristics,  
    _In_      BOOLEAN Exclusive,  
    _Outptr_  PDEVICE_OBJECT *DeviceObject);
```


Device Symbolic Link

- Drivers should define a second name, if the device object is created with a name
- The kernel device name is not visible from user mode (**CreateFile**)
 - However, it is possible to access with the native **NtOpenFile**
 - Must create a symbolic link to the kernel device object
 - The name should be under the "\\??\\" directory in the object manager's namespace
 - The suffix of both names need not be the same, but usually is

Creating a Symbolic Link

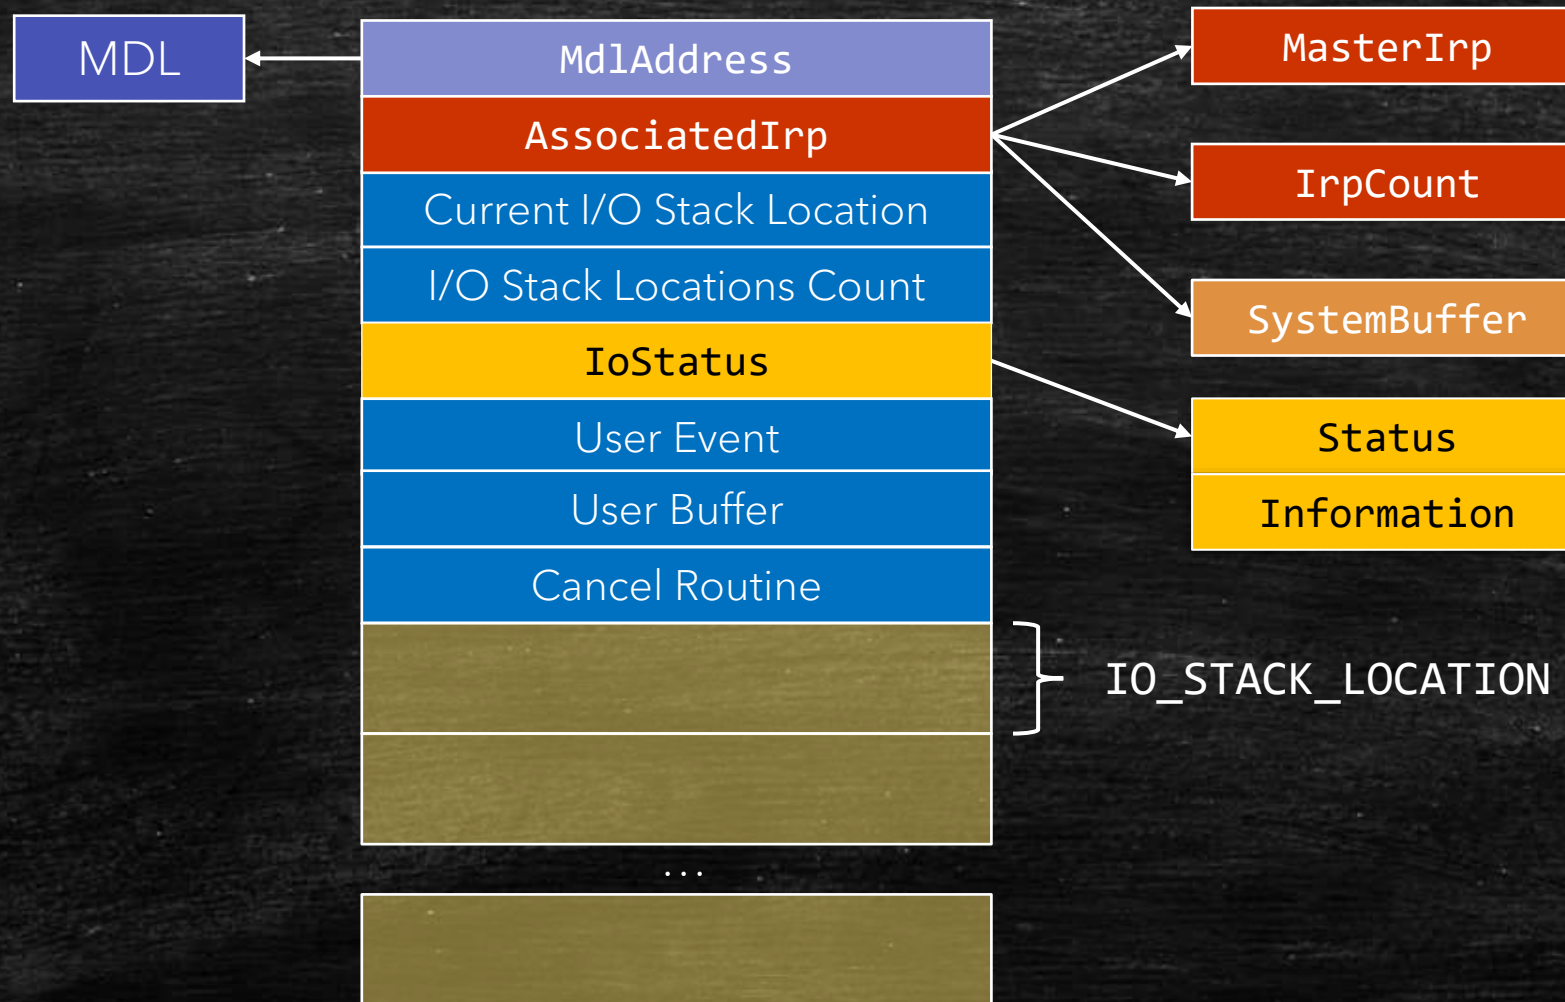
- Call `IoCreateSymbolicLink`
- "\\??\\" is a subdirectory in the object manager's namespace
 - Run *WinObj* tool from SysInternals to view
 - User mode apps can enumerate with `QueryDosDevice`

```
UNICODE_STRING Win32DevName;  
RtlInitUnicodeString(&Win32DevName, L"\\??\\MyDev1");  
IoCreateSymbolicLink(&Win32DevName, &DeviceName);
```


I/O Request Packet (IRP)

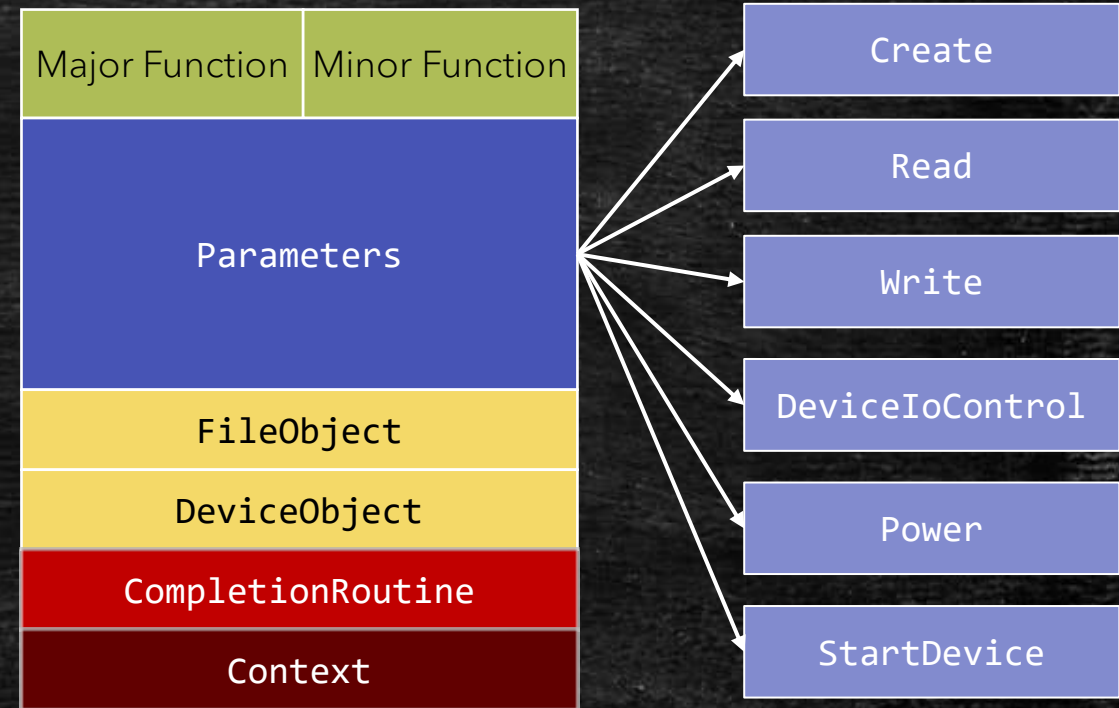
- A structure representing some request
 - Represented by the IRP structure
 - Contains all details needed to handle the request (codes, buffers, sizes, etc.)
- Always allocated from non-paged pool
- Accompanied by a set of structures of type `IO_STACK_LOCATION`
 - Complements the data in the IRP
- IRPs are typically created by the I/O Manager, P&P Manager or the Power Manager

The IRP Structure



The I/O Stack Location

- Each created IRP comes with an array of **IO_STACK_LOCATION** structures
- Basically, contain more info about the request



Dispatch Routines

- Located through the Driver object's **MajorFunction** array
 - Selected by the IRP major function code (located in the current I/O stack location)
- An unassigned entry means “unimplemented function”
 - I/O Manager returns error automatically

```
NTSTATUS SomeDispatchRoutine(  
    _In_ PDEVICE_OBJECT DeviceObject,  
    _In_ PIRP Irp);
```


Common IRP Types

- **IRP_MJ_CREATE**
 - Invoked for **CreateFile** calls
- **IRP_MJ_CLOSE**
 - Invoked when last handle closed to the file object
- **IRP_MJ_READ, IRP_MJ_WRITE**
 - Invoked for **ReadFile, WriteFile**
- **IRP_MJ_DEVICE_CONTROL**
 - Invoked for **DeviceIoControl**

Completing the Request

```
Irp->IoStatus.Status = STATUS_XXX;  
Irp->IoStatus.Information = NumberOfBytesTransferred;  
IoCompleteRequest(Irp, IO_NO_INCREMENT);  
return STATUS_XXX;
```

- Returned status is mapped to a Win32 error code
 - There is no one to one mapping relationship
 - See *NTSTATUS.H* for possible error codes
- **NumberOfBytesTransferred** should be zero if error occurred
 - Might mean something else depending on the type of IRP
- Do not boost thread priority on immediate completion

Installing a Software Driver

- Normally, an INF file is not needed
- Use the **CreateService** API to install
 - Or a comparable tool, such as **Sc.exe**
- Example
 - Use elevated command window
 - Mind the spaces (or lack thereof)

```
sc.exe create MyDriver type= kernel binPath= c:\Source\MyDriver.sys
```

Starting and Stopping the Driver

- By default, driver is installed with start type "on demand"
- To load the driver, use the **StartService** API
 - Or "`sc start mydriver`"
- Driver binary is loaded into kernel space
- **DriverEntry** is now called
- To stop the driver
 - Call the **ControlService** API with **SERVICE_CONTROL_STOP**
 - Or use "`sc stop mydriver`"
 - Unload routine is called before the driver binary is unloaded

Testing the Driver

- Copy the binary (SYS) file to the target system
 - Preferably a virtual machine
- Enable test signing mode so that unsigned drivers can be loaded
 - `bcdedit /set testsigning on` (and restart)
- Install the driver on the target system
- Copy a test application to the target system
- Use the test application to send commands to the driver

Debugging the Driver

- Install the driver on the target system
- Create a debugger connection to the target system
- Make sure the debugger symbol path points to where the PDB file is
 - Or copy the PDB file to the same location as the driver
- Start remote kernel debugging
- Set breakpoints as appropriate
 - Use the **bu** command to set a breakpoint at **DriverEntry**
- Debug!

KdPrint / DbgPrint Output

- Using the **KdPrint** macro or **DbgPrint** function sends output to a debugger, if connected
- Requires creating a registry key *HKLM\System\CurrentControlSet\Control\Session Manager\Debug Print Filter*
 - Add a DWORD value named **DEFAULT** and set its value to 8
 - Restart the system
- Can capture debug output with the *DbgView* Sysinternals tool
 - Use the option "Enable Verbose Kernel Output" to capture calls regardless of the above Registry key/value settings

Remote Kernel Debugging

- Target machine
 - Configure for debugging as before
 - Select a communication medium
 - *MsConfig.exe* or *bcdedit.exe /dbgsettings*
 - Serial, USB, Network (Windows 8+)
- Host machine
 - File | Kernel Debug...
 - Select configured communication medium
- If target is a virtual machine
 - Can expose a VM COM port as a host named pipe
 - Or use network if host and target are Windows 8+

Remote Kernel Debugging (Network)

- Target
 - Elevated command window
 - `bcdedit /dbgsettings net hostip:<ip> port:<port> [key:<key>]`
 - If key is not provided, a random key is generated
 - Restart the system
- Host (debugger)
 - Select network debugging
 - Set same port and key
- The *KdNet.exe* tool simplifies the above procedure somewhat